

# Course Project Writeup

Richard Creamer

Friday, August 22, 2014

## Summary

### Data Preparation

- I chose to use a Random Forest
- Since the implementation of Random Forest handles cross-validation, I only created Train and Test sets (80:20)

### Model Selection

- I followed the lecture's advice and used `rfcv()` to evaluate cross-validation OOB error estimates
- I subsetting data down to 53 columns including 'classe' y column (discard column if NA/blank count > nrow/2)
- I used ranked Linear Model coefficients to extract 20 features which should perform well
- I also used the Random Forest `importance()` output's feature error estimates for possible feature reduction, but rejected this because the plots (below) do not clearly delineate less-important features
- I then fit Random Forests over a grid/range of 'mtry' and 'ntree' parameters to find the best-performing values for these parameters for the full 52-feature training set (rf52)
- I did the same thing (grid/range of 'mtry' and 'ntree') for a 20-features derived data frame using the Linear Model coefficients for subsetting

### Model Evaluation and OOB Error Estimates

- I then evaluated the performance of rf52 and rf20 on the TEST set
- I then compared the TEST set performance for the rf52 and rf20 Random Forests to the `rfcv()` and the Random Forests' built-in OOB estimates:

#### Random Forest rf52 (trained using 52 features, mtry=2, ntree=20)

- Train accuracy: 1.00000
- TEST set prediction accuracy: 99.1847%
- Random Forest rf52 built-in OOB error estimate: 2.3445%

#### Random Forest rf20 (trained using 20 features, mtry=2, ntree=50)

- Train accuracy: 1.00000
- TEST set prediction accuracy: 99.4395%
- Random Forest rf20 built-in OOB error estimate: 0.91%

#### rfcv() OOB Error Estimates by Variable Count

- 52 variables: 0.005542460
- 26 variables: 0.007581066
- 13 variables: 0.009683379
- 6 variables: 0.045104160

- 3 variables: 0.109320252
- 1 variable: 0.596929350

### Prediction on Project 20-row Prediction Set (pml-testing.csv)

- rf52: B A B A A E D B A A B C B A E E A B B B
- rf20: B A B A A E D B A A B C B A E E A B B B

### Discussion

The rf20 Random Forest had slightly better performance than the rf52 RF even though the rf20 RF was trained using only 20 features. I think the reason for this is that my grid search parameter optimization function assigned 50 trees to the rf20 RF whereas it only assigned 20 trees to the rf52 RF.

### Notes

- I added a separate PDF file containing this R Markdown Course Project write-up
- I added a separate PDF file to my repo containing the test script's output (plots manually inserted)
- I added a separate CpScript.R file to my repo containing the test script 'driver' and various other 'helper' functions
- The R code is rather messy do to the sequence of steps taken and the print statements (sorry)
- The CpScript.R has a 'driver' test function named cpTest(), but only the steps inside this function are listed in this Rmd write-up, not the outer function itself
- Many of my helper functions return tuples containing multiple named return values

IMPORTANT: I could not get knitr to work with code chunks which call functions in another chunk. Therefore, I am having to submit the code in chunks but NOT labeled as {r} code, and instead, am including a code chunk at the bottom with my entire (large) .R file. As a result, the chunks in the write-up do not have syntax highlighting/coloring. This also forced me to manually insert both the code's output as well as HTML tags to intersperse my plots' PNG files instead of being generated by R/knitr. SORRY! Full code at bottom WITH syntax highlighting.

## Begin 'Test Driver' Code Chunks, Output, and Plots

### Loading and Processing the Data

Note: return value 'd' is a tuple

Code:

```
# load data, create and split data frames, shuffle rows, results stored in d tuple
d <- loadData( subDir, doPrint = TRUE )
```

Output:

```
[1] ---> loading data...
[1] ...done reading data and creating data frames
```

### Assessing rfcv() Error Data for Feature Selection

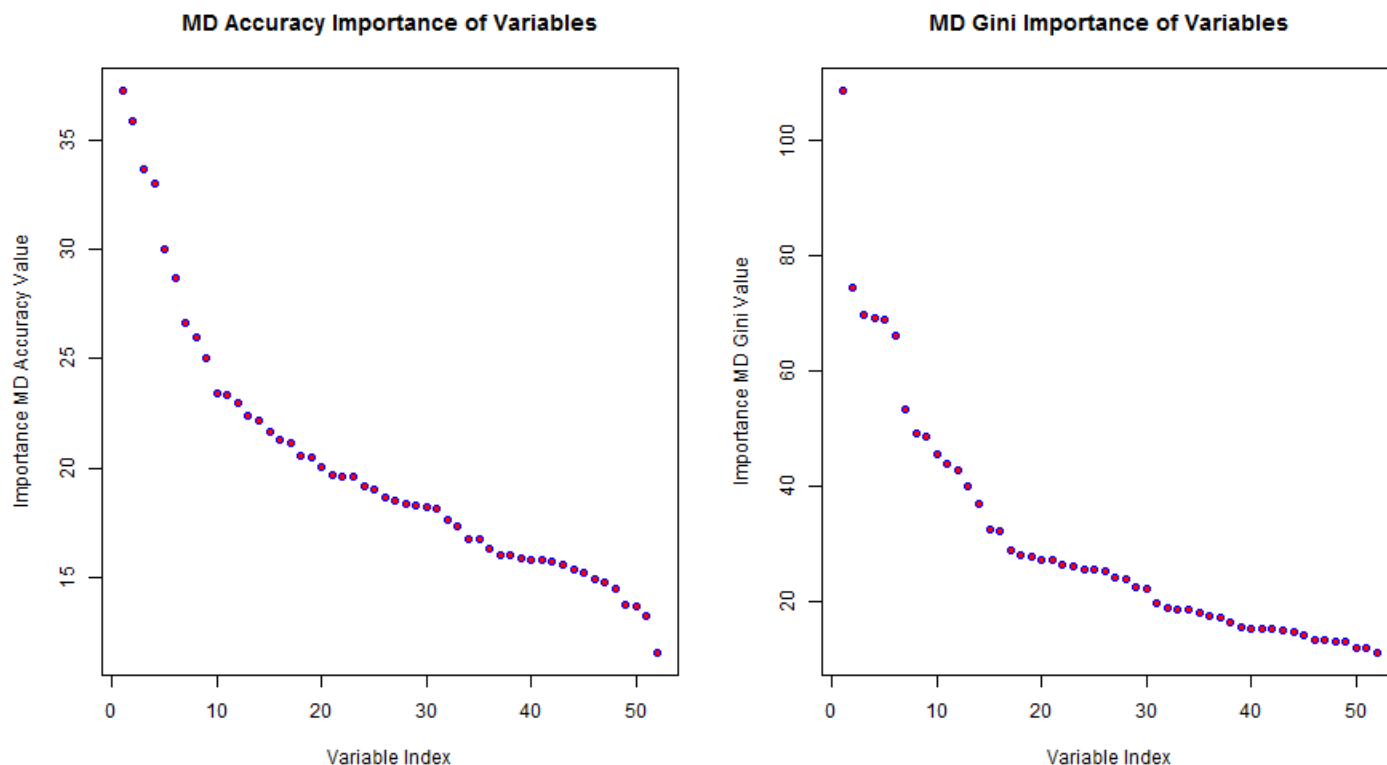
Note: the plot resulting from the code, below, did not contain any obvious pattern to use for separating important from less-important features.

Code:

```
# see if rfcv() variable importance useful for features selection
evalRfcvVarImport( d$trainDf, nrows=2000, doPrint=TRUE )
```

Output:

```
[1] plotting Random Forest variable importance() metrics  
[1] no clear dividing line differentiating important vs. unimportant variables
```



### Fitting a Linear Model and Extracting 20 Largest Coefficients


Note: the plot resulting from the code below DID show that the error flattened out after using approximately 20 features.

Code:

```
# get sorted coefficients from Linear Model fit to see if useful  
# for feature selection  
lmBest20Coeffs <- getMostImportantLmCoeffs( d$trainDf, nrows=1000, doPrint=TRUE )
```

Output:

```
[1] ---> assessing whether Linear Model coefficients offer useful variable importance  
rankings...  
[1] plotting Linear Model coefficients sorted in decreasing order
```



### Finding Best 'ntree' and 'mtry' Random Forests for 52-feature TRAIN Set

Note: below, the code calls a function which chooses the best Random Forest and its 'mtry' and 'ntree' parameters over a range of values for these two variables.

Code:

```
# fit Random Forests over a grid of ranges for params 'mtry' and 'ntree'
# use full 52 features in training set
pr( "---> examining Random Forest performance for a range of 'mtry' and 'ntree'
parameters..." )
mtryVals <- c( 2, 5, 10 )
ntreeVals <- c( 1, 2, 3, 5, 10, 20 )
rf52 <- findBestRfParams( d$trainDf, mtryVals, ntreeVals, doPrint = TRUE ) # rf52 is a
tuple
pr( "---> Random Forest accuracy values for evaluated (mtry, ntree) grid pairs:" )
pr( "---> NOTE: row names are mtry values; column names are ntree values" )
pr( "---> NOTE: matrix entries are classification accuracy on train set" )
cat( "\n" )
pr( rf52$accMatrix )
cat( "\n" )
pr( "--->parameters from best 52-feature Random Forest:" )
pr( sprintf( "resultant rf52: train accuracy=%f mtry=%d ntree=%d OOB error=%f",
  rf52$bestAcc, rf52$bestMtry, rf52$bestNtree, getOob( rf52$bestRf ) ) )
cat( "\n" )
```

```
# now just print out the Random Forest to get confusion matrix and OOB error est.
pr( "---> printing best (rf52) Random Forest for above parameters: " )
pr( rf52$bestRf )
cat( "\n" )
```

Output:

```
[1] ---> examining Random Forest performance for a range of 'mtry' and 'ntree' parameters...
```

```
[1] evaluating Random Forest w/ mtry=2 ntree=1
[1] evaluating Random Forest w/ mtry=2 ntree=2
[1] evaluating Random Forest w/ mtry=2 ntree=3
[1] evaluating Random Forest w/ mtry=2 ntree=5
[1] evaluating Random Forest w/ mtry=2 ntree=10
[1] evaluating Random Forest w/ mtry=2 ntree=20
[1] evaluating Random Forest w/ mtry=5 ntree=1
[1] evaluating Random Forest w/ mtry=5 ntree=2
[1] evaluating Random Forest w/ mtry=5 ntree=3
[1] evaluating Random Forest w/ mtry=5 ntree=5
[1] evaluating Random Forest w/ mtry=5 ntree=10
[1] evaluating Random Forest w/ mtry=5 ntree=20
[1] evaluating Random Forest w/ mtry=10 ntree=1
[1] evaluating Random Forest w/ mtry=10 ntree=2
[1] evaluating Random Forest w/ mtry=10 ntree=3
[1] evaluating Random Forest w/ mtry=10 ntree=5
[1] evaluating Random Forest w/ mtry=10 ntree=10
[1] evaluating Random Forest w/ mtry=10 ntree=20
```

```
[1] --> Random Forest accuracy values for evaluated (mtry, ntree) grid pairs:
```

```
[1] --> NOTE: row names are mtry values; column names are ntree values
```

```
[1] --> NOTE: matrix entries are classification accuracy on train set
```

	1	2	3	5	10	20
2	0.9548321	0.9512009	0.9906989	0.9978977	0.9997452	1.0000000
5	0.9720966	0.9693572	0.9950946	0.9984710	0.9998726	1.0000000
10	0.9684016	0.9678282	0.9948398	0.9984073	0.9998726	0.9999363

```
[1] --->parameters from best 52-feature Random Forest:
```

```
[1] resultant rf52: train accuracy=1.000000 mtry=2 ntree=20 OOB error=0.023445
```

```
[1] ---> printing best (rf52) Random Forest for above parameters:
```

Call:

```
randomForest(formula = classe ~ ., data = df, mtry = mt, ntree = nt)
```

```
  Type of random forest: classification
```

```
    Number of trees: 20
```

```
No. of variables tried at each split: 2
```

```
  OOB estimate of  error rate: 2.34%
```

Confusion matrix:

	A	B	C	D	E	class.error
A	4394	22	9	14	2	0.01058320
B	45	2951	32	14	10	0.03309305
C	7	45	2610	30	7	0.03297518
D	12	3	63	2480	10	0.03426791
E	2	15	4	22	2893	0.01464578

## Using rfcv() to Evaluate CV OOB Error

Code:

```
# now use rfcv() to evaluate cross-validation error
pr( "---> computing rfcv() cross-validation error - this may take several minutes..." )
set.seed( 1 )
rfcvOutput <- rfcv( d$trainDf[, -53], d$trainDf[, 53] )
pr( "...done computing rfcv() output" )
pr( "rfcv() cross-validation estimates for training set vs. number variables used:")
pr( rfcvOutput$error.cv )
cat( "\n" )
```

Output:

```
[1] ---> computing rfcv() cross-validation error - this may take several minutes...
[1] ...done computing rfcv() output
[1] rfcv() cross-validation estimates for training set vs. number variables used:
      52      26      13      6      3      1
0.005542460 0.007581066 0.009683379 0.045104160 0.109320252 0.596929350
```

## Evaluating TEST set accuracy of Random Forest (rf52) Trained on 52 Features

Code:

```
# compute performance of rf52 best 52-feature random forest on T-E-S-T set
acc <- evalRf( rf52$bestRf, d$testDf )
pr( sprintf( "accuracy of best (rf52) Random Forest on 20%% TEST set: %f", acc ) )
cat( "\n" )
```

Output:

```
[1] accuracy of best (rf52) Random Forest on 20% TEST set: 0.991847
```

## Finding Best 'ntree' and 'mtry' Random Forests for 20-feature TRAIN Set

Note: because I was using fewer features, the 'mtry' and 'ntree' ranges I used for evaluation ranged to higher values than for the rf52 Random Forest grid search.

Code:

```
# additional exercise: find best 20-feature Random Forest using Linear Model top-20
coeffs
pr( "---> As additional exercise fit Random Forest to top-20 features from Linear Model" )
lmBest20Coeffs[ 21 ] = "classe" # Need to append the 'classe' to feature list
best20TrainDf <- d$trainDf[, lmBest20Coeffs ] # subset training set, only the 20 'top'
features
mtryVals <- c( 2, 5, 10 ) # use different parameter ranges as expect lower accuracy using
fewer features
ntreeVals <- c( 10, 50, 100, 200 ) # ditto
rf20 <- findBestRfParams( d$trainDf, mtryVals, ntreeVals, doPrint = TRUE )
pr( "--> Random Forest accuracy values for evaluated (mtry, ntree) grid pairs:" )
pr( "--> NOTE: row names are mtry values; column names are ntree values" )
pr( "--> NOTE: matrix entries are classification accuracy on train set" )
cat( "\n" )

pr( rf20$accMatrix )
cat( "\n" )
pr( "--->parameters from best 20-feature Random Forest:" )
pr( sprintf( "resultant rf20: train accuracy=%f mtry=%d ntree=%d OOB error=%f",
```

```

      rf20$bestAcc, rf20$bestMtry, rf20$bestNtree, getOob( rf20$bestRf ) ) )
cat( "\n" )

# PRINT BEST RANDOM FOREST (to get OOB and confusion matrix)
pr( "---> printing best Random Forest for above parameters: " )
pr( rf20$bestRf )
cat( "\n" )

```

Output:

```

[1] ---> As additional exercise fit Random Forest to top-20 features from Linear Model
[1] evaluating Random Forest w/ mtry=2 ntree=10
[1] evaluating Random Forest w/ mtry=2 ntree=50
[1] evaluating Random Forest w/ mtry=2 ntree=100
[1] evaluating Random Forest w/ mtry=2 ntree=200
[1] evaluating Random Forest w/ mtry=5 ntree=10
[1] evaluating Random Forest w/ mtry=5 ntree=50
[1] evaluating Random Forest w/ mtry=5 ntree=100
[1] evaluating Random Forest w/ mtry=5 ntree=200
[1] evaluating Random Forest w/ mtry=10 ntree=10
[1] evaluating Random Forest w/ mtry=10 ntree=50
[1] evaluating Random Forest w/ mtry=10 ntree=100
[1] evaluating Random Forest w/ mtry=10 ntree=200

[1] --> Random Forest accuracy values for evaluated (mtry, ntree) grid pairs:
[1] --> NOTE: row names are mtry values; column names are ntree values
[1] --> NOTE: matrix entries are classification accuracy on train set

```

		10	50	100	200
2	0.9997452	1	1	1	1
5	0.9998726	1	1	1	1
10	0.9998726	1	1	1	1

```

[1] --->parameters from best 20-feature Random Forest:
[1] resultant rf20: train accuracy=1.000000 mtry=2 ntree=50 OOB error=0.009110

[1] ---> printing best Random Forest for above parameters:

```

Call:

```

randomForest(formula = classe ~ ., data = df, mtry = mt, ntree = nt)
      Type of random forest: classification
      Number of trees: 50
No. of variables tried at each split: 2

```

OOB estimate of error rate: 0.91%

Confusion matrix:

	A	B	C	D	E	class.error
A	4430	7	0	3	1	0.002476920
B	21	3020	8	1	2	0.010484928
C	3	28	2663	5	1	0.013703704
D	3	0	42	2519	4	0.019080997
E	0	3	1	10	2922	0.004768392

## Evaluating TEST set accuracy of Random Forest (rf20) Trained on 20 Features

Code:

```
# compute performance of rf20 Random Forest on T-E-S-T set (20% of train set rows)
pr( "---> evaluate top-20 feature Random Forest: " )
acc <- evalRf( rf20$bestRf, d$testDf )
pr( sprintf( "accuracy of best (rf20) Random Forest on 20%% TEST set: %f", acc ) )
cat( "\n" )
```

Output:

```
[1] ---> evaluate top-20 feature Random Forest:
[1] accuracy of best (rf20) Random Forest on 20% TEST set: 0.994395
```

## Using rf52 and rf20 Models to Predict Labels for 20-Row Project Data (pml-testing.csv)

Code:

```
# use rf52 to predict labels for course project 20-row data set (pml-testing.csv)
pr( "---> use rf52 to predict labels for course project 20-row data set (pml-testing.csv"
)
predsRf52 <- predict( rf52$bestRf, d$predDf )
pr( predsRf52 )
cat( "\n" )

# use rf20 to predict labels for course project 20-row data set (pml-testing.csv)
pr( "---> use rf20 to predict labels for course project 20-row data set (pml-testing.csv"
)
predsRf20 <- predict( rf20$bestRf, d$predDf )
pr( predsRf20 )
cat( "\n" )
```

Output:

```
[1] ---> use rf52 to predict labels for course project 20-row data set (pml-testing.csv
 1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20
  B  A  B  A  A  E  D  B  A  A  B  C  B  A  E  E  A  B  B  B
Levels: A B C D E

[1] ---> use rf20 to predict labels for course project 20-row data set (pml-testing.csv
 1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20
  B  A  B  A  A  E  D  B  A  A  B  C  B  A  E  E  A  B  B  B
Levels: A B C D E
```

## Full R Code with Syntax Highlighting and Coloring

```
#####
# Usage:
# source( "CpScript.R" )
# set 'subDir' variable if != local sub-directory 'data'
# cpTest()
#####

library( randomForest )

## Warning: package 'randomForest' was built under R version 3.1.1

## randomForest 4.6-10
## Type rfNews() to see new features/changes/bug fixes.
```



```

cpScript <- function( subDir = "data" ) {

  repoDir <-
"D:/OldLaptop/D_Drive/JohnsHopkinsCoursera/8_PracticalMLCoursera/hw/CourseProjRepo"

  # Load data, create and split data frames, shuffle rows, results stored in d tuple
  d <- loadData( subDir, doPrint = TRUE )

  # see if rfcv() variable importance useful for features selection
  evalRfcvVarImport( d$trainDf, nrows=2000, doPrint=TRUE )

  # get sorted coefficients from Linear Model fit to see if useful
  # for feature selection
  lmBest20Coeffs <- getMostImportantLmCoeffs( d$trainDf, nrows=1000, doPrint=TRUE )

  # fit Random Forests over a grid of ranges for params 'mtry' and 'ntree'
  # use full 52 features in training set
  pr( "---> examining Random Forest performance for a range of 'mtry' and 'ntree'
parameters..." )
  mtryVals <- c( 2, 5, 10 )
  ntreeVals <- c( 1, 2, 3, 5, 10, 20 )
  rf52 <- findBestRfParams( d$trainDf, mtryVals, ntreeVals, doPrint = TRUE ) # rf52 is
a tuple
  pr( "--> Random Forest accuracy values for evaluated (mtry, ntree) grid pairs:" )
  pr( "--> NOTE: row names are mtry values; column names are ntree values" )
  pr( "--> NOTE: matrix entries are classification accuracy on train set" )
  cat( "\n" )
  pr( rf52$accMatrix )
  cat( "\n" )
  pr( "--->parameters from best 52-feature Random Forest:" )
  pr( sprintf( "resultant rf52: train accuracy=%f mtry=%d ntree=%d OOB error=%f",
    rf52$bestAcc, rf52$bestMtry, rf52$bestNtree, getOob( rf52$bestRf ) ) )
  cat( "\n" )

  # now just print out the Random Forest to get confusion matrix and OOB error est.
  pr( "---> printing best (rf52) Random Forest for above parameters: " )
  pr( rf52$bestRf )
  cat( "\n" )

  # now use rfcv() to evaluate cross-validation error
  pr( "---> computing rfcv() cross-validation error - this may take several minutes..."
)

  set.seed( 1 )
  rfcvOutput <- rfcv( d$trainDf[, -53], d$trainDf[, 53] )
  pr( "...done computing rfcv() output" )
  pr( "rfcv() cross-validation estimates for training set vs. number variables used:")
  pr( rfcvOutput$error.cv )
  cat( "\n" )

  # compute performance of rf52 best 52-feature random forest on T-E-S-T set
  acc <- evalRf( rf52$bestRf, d$testDf )
  pr( sprintf( "accuracy of best (rf52) Random Forest on 20%% TEST set: %f", acc ) )
  cat( "\n" )

  # additional exercise: find best 20-feature Random Forest using Linear Model top-20

```

```

coeffs
  pr( "---> As additional exercise fit Random Forest to top-20 features from Linear
Model" )
  lmBest20Coeffs[ 21 ] = "classe" # Need to append the 'classe' to feature List
  best20TrainDf <- d$trainDf[, lmBest20Coeffs ] # subset training set, only the 20
'top' features
  mtryVals <- c( 2, 5, 10 ) # use different parameter ranges as expect Lower accuracy
using fewer features
  ntreeVals <- c( 10, 50, 100, 200 ) # ditto
  rf20 <- findBestRfParams( d$trainDf, mtryVals, ntreeVals, doPrint = TRUE )
  pr( "--> Random Forest accuracy values for evaluated (mtry, ntree) grid pairs:" )
  pr( "--> NOTE: row names are mtry values; column names are ntree values" )
  pr( "--> NOTE: matrix entries are classification accuracy on train set" )
  cat( "\n" )

  pr( rf20$accMatrix )
  cat( "\n" )
  pr( "--->parameters from best 20-feature Random Forest:" )
  pr( sprintf( "resultant rf20: train accuracy=%f mtry=%d ntree=%d OOB error=%f",
    rf20$bestAcc, rf20$bestMtry, rf20$bestNtree, get0ob( rf20$bestRf ) ) )
  cat( "\n" )

# PRINT BEST RANDOM FOREST (to get OOB and confusion matrix)
pr( "---> printing best Random Forest for above parameters: " )
pr( rf20$bestRf )
cat( "\n" )

# compute performance of rf20 Random Forest on T-E-S-T set (20% of train set rows)
pr( "---> evaluate top-20 feature Random Forest: " )
acc <- evalRf( rf20$bestRf, d$testDf )
pr( sprintf( "accuracy of best (rf20) Random Forest on 20%% TEST set: %f", acc ) )
cat( "\n" )

# use rf52 to predict labels for course project 20-row data set (pml-testing.csv)
pr( "---> use rf52 to predict labels for course project 20-row data set (pml-
testing.csv" )
predsRf52 <- predict( rf52$bestRf, d$predDf )
pr( predsRf52 )
cat( "\n" )

# use rf20 to predict labels for course project 20-row data set (pml-testing.csv)
pr( "---> use rf20 to predict labels for course project 20-row data set (pml-
testing.csv" )
predsRf20 <- predict( rf20$bestRf, d$predDf )
pr( predsRf20 )
cat( "\n" )

# RETURN LIST OF SELECTED VARIABLES FOR CALLER
list( data=d, lmBest20Coeffs=lmBest20Coeffs, rf52=rf52, rf20=rf20,
  rfcvOutput=rfcvOutput, predsrf52=predsr52, predsRf20=predsRf20 )
}

```

```
#####
```

```

# helper function which fits a Linear Model and returns
# return value = 20 largest sorted coefficients of Linear Model

```

```

getMostImportantLmCoeffs <- function( df, nrows=0, doPrint=FALSE ) {
  if ( nrows == 0 )
    nrows = nrow( df )
  # EVAL LINEAR MODEL COEFFICIENT RANKING FOR *FEATURE SELECTION* (result: useful)
  # now try ranking variables by Linear Model coefficient values
  if ( doPrint )
    pr( "---> assessing whether Linear Model coefficients offer useful variable
importance rankings..." )
  classeIdx <- getColIdx( df, "classe" )
  y <- as.numeric( df[ 1:nrows, classeIdx ] ) # use only 1,000 rows for LM fit
  lmDf <- cbind( df[ 1:nrows , 1:classeIdx - 1 ], y )
  set.seed( 1 ) # set RNG seed for reproducibility
  lmFit <- lm( y ~ . , data = lmDf )
  sortedCoeffs <- sort( abs( lmFit$coefficients ), decreasing = TRUE )
  nCoeffs <- length( sortedCoeffs )
  sortedCoeffNames <- names( sortedCoeffs[2:nCoeffs ] )
  lm20MostImpFeatures <- names(sortedCoeffs)[2:21] # skip intercept = coef[1]
  # plot linear model coefficients largest-to-smallest; skipping intercept coefficient
  plotLmCoeffVals( 2:nCoeffs, sortedCoeffs[2:nCoeffs], doPrint=doPrint )
  if ( doPrint )
    cat( "\n" )
  lmBest20Coeffs <- sortedCoeffNames[2:21]
}

#####
# helper function to use rfcv() output ranking of variables to see
# if the rankings are useful for feature selection/reduction
evalRfcvVarImport <- function( df, nrows=0, doPrint = FALSE ) {
  # evaluate rfcv() output to see if useful for feature selection
  if ( nrows == 0 )
    nrows = nrow( df )
  set.seed( 1 )
  rf <- randomForest( classe ~ . , data = df[1:nrows, ], importance = TRUE )
  impVal <- as.data.frame( importance( rf ) )
  impValMda <- impVal[ rev( order( impVal$MeanDecreaseAccuracy ) ), ]
  impValGini <- impVal[ rev( order( impVal$MeanDecreaseGini ) ), ]
  mdaDf <- data.frame( rownames( impValMda ), impValMda$MeanDecreaseAccuracy )
  giniDf <- data.frame( rownames( impValGini ), impValGini$MeanDecreaseGini )
  nVars <- nrow( mdaDf )
  # plot the results
  par( mfrow = c( 1, 2 ) )
  plotImportanceData( 1:nVars, mdaDf[,2], giniDf[,2], doPrint=doPrint )
  # plot conclusion - no clear dividing line between important vs. unimportant
variables
  # results not actionable
  if ( doPrint ) {
    pr( "no clear dividing line differentiating important vs. unimportant variables"
)
    cat( "\n" )
  }
}

#####
# helper function to load, subset, and shuffle data
# output: various data frames

```

```

# read and subset/process 2 CSV files, create data frames
loadData <- function( subDir = "data", doPrint = FALSE ) {
  if ( doPrint )
    pr( "---> loading data..." )
  trainFile <- file.path( subDir, "pml-training.csv" )
  testFile <- file.path( subDir, "pml-testing.csv" )
  trainFileDf <- prepDf( read.csv( trainFile ) )
  predDf <- prepDf( read.csv( testFile ) ) # 20-row prediction file
  set.seed( 1 ) # set seed for shuffle operation
  nrows <- nrow( trainFileDf )
  trainFileDf <- trainFileDf[ sample( nrows ), ] # randomly shuffle rows
  # Partition: 80% train, 20% test
  trainDf <- trainFileDf[ 1 : as.integer( 0.8 * nrows ), ]
  testDf <- trainFileDf[ as.integer( nrow( trainDf ) + 1 ) : nrows, ]
  if ( doPrint ) {
    pr( "...done reading data and creating data frames" )
    cat( "\n" )
  }
  list( trainDf=trainDf, testDf=testDf, predDf=predDf )
}

#####
# helper function to evaluate Random Forests over range of mtry and ntree parameter
findBestRfParams <- function( df, mtryVals, ntreeVals, doPrint = FALSE ) {
  bestAcc <- 0
  bestMtry <- 0
  bestNtree <- 0
  bestRf <- NULL
  accMatrix <- matrix( nrow=length( mtryVals ), ncol=length( ntreeVals ) )
  rownames( accMatrix ) <- as.character( mtryVals )
  colnames( accMatrix ) <- as.character( ntreeVals )
  for ( i in 1:length( mtryVals ) ) { # mtry
    mt <- mtryVals[ i ]
    for ( j in 1:length( ntreeVals ) ) { # ntree
      nt <- ntreeVals[ j ]
      if ( doPrint )
        pr( sprintf( "evaluating Random Forest w/ mtry=%d ntree=%d", mt, nt ) )
      set.seed( 1 )
      rf <- randomForest( classe ~ ., data=df, mtry=mt, ntree=nt )
      acc <- evalRf( rf, df )
      accMatrix[i, j] = acc
      if ( acc > bestAcc ) {
        bestMtry <- mt
        bestNtree <- nt
        bestRf <- rf
        bestAcc <- acc
      }
    }
  }
  if ( doPrint )
    cat( "\n" )
  list( bestAcc=bestAcc, bestMtry=bestMtry, bestNtree=bestNtree,
        bestRf=bestRf, accMatrix=accMatrix )
}

```

```

#####
# helper method to get OOB error estimate from Random Forest
getOob <- function( rf ) {
  rf$err.rate[rf$ntree, 1 ]
}

#####
# helper method to evaluate classification accuracy of Random Forest wrt a data frame
evalRf <- function( rf, df ) {
  predTestSet <- predict( rf, df ) # use Test set, NOT Cross-Validation set
  numAgree <- sum( predTestSet == df$classe )
  modelTestAccur <- numAgree/length( predTestSet )
}

#####
# helper method to reduce line lengths
pr <- function( msg ) {
  print( msg, quote = FALSE )
}

#####
# helper to de-clutter script code
plotLmCoeffVals <- function( coefIdxs, coefVals, doPrint=FALSE ) {

  if ( doPrint )
    pr( "plotting Linear Model coefficients sorted in decreasing order" )

  # plot to screen and knitr
  plot( coefIdxs, coefVals, pch=21, col="blue", bg="red",
        xlab="Linear Model Coef Index",
        ylab="LM Coef Value",
        main="Linear Model Coefficients" )
  lines(coefIdxs, coefVals, col="blue" )

  # also plot to PNG file
  png( "LmCoeff.png", height = 512, width = 512 )
  par( family = "sans" )
  plot( coefIdxs, coefVals, pch=21, col="blue", bg="red",
        xlab="Linear Model Coef Index",
        ylab="LM Coef Value",
        main="Linear Model Coefficients" )
  lines(coefIdxs, coefVals, col="blue" )
  dev.off()

  if ( doPrint )
    cat( "\n" )
}

#####
# helper to de-clutter script code
plotImportanceData <- function( varIndices, mdaVals, mdGiniVals, doPrint=FALSE ) {

  if ( doPrint )

```

```

    pr( "plotting Random Forest variable importance() metrics" )

# plot to screen and knitr
plot( varIndices, mdaVals,
      xlab="Variable Index",
      ylab="Importance MD Accuracy Value",
      main="MD Accuracy Importance of Variables",
      pch=21, col="blue", bg="red" )
plot( varIndices, mdGiniVals,
      xlab="Variable Index",
      ylab="Importance MD Gini Value",
      main="MD Gini Importance of Variables",
      pch=21, col="blue", bg="red" )

# plot to file as well
png( "RfImpVarMetrics.png", height = 512, width = 900 )
par( family = "sans" )
par( mfrow = c( 1, 2 ) )
plot( varIndices, mdaVals,
      xlab="Variable Index",
      ylab="Importance MD Accuracy Value",
      main="MD Accuracy Importance of Variables",
      pch=21, col="blue", bg="red" )
plot( varIndices, mdGiniVals,
      xlab="Variable Index",
      ylab="Importance MD Gini Value",
      main="MD Gini Importance of Variables",
      pch=21, col="blue", bg="red" )
dev.off()
if ( doPrint )
  cat( "\n" )
}

#####
# helper function to subset, coerce data
prepDf <- function( df ) {
  ## discard first 7 columns - may not be good for general data sets
  df <- df[ , 8:ncol( df ) ]
  ## discard columns with number na/blanks > num rows in data frame
  colNaSums <- apply( df, 2, function(x) { length( which ( is.na(x) | x == "" ) ) } )
  df <- df[ , colNaSums < nrow(df)/2 ]
  ## coerce integers to numeric
  for ( i in 1:ncol( df ) ) {
    if ( class( df[ 1, i ] ) == "integer" )
      df[ , i ] <- as.numeric( df[ , i ] )
  }
  df # Caller must do the shuffling
}

#####
# a function which returns the numeric index of column in a data frame given the column name
getColIdx <- function( df, colName ) {
  grep( colName, colnames( df ) )
}

```